

JOS-Lab 5: File System

Disk Access

Exercise 1

Modify `env_create()` to give I/O privileges to the ring3 Env which type is `ENV_TYPE_FS` by setting the `IOPL` to 3 in the `eflags`.

```
if (type == ENV_TYPE_FS) {
    e->env_tf.tf_eflags |= FL_IOPL_3;
}
```

The Block Cache

Exercise 2

In `bc_pgfault()`, for a page fault within the block cache region, it first allocate a page for that page:

```
void *pgaddr = ROUNDDOWN(addr, PGSIZE);
r = sys_page_alloc(0, pgaddr, PTE_W|PTE_U|PTE_P);
if (r < 0)
    panic("sys_page_alloc: %e", r);
```

Then load the contents of the block from the disk into that page:

```
r = ide_read(blockno * BLKSECTS, pgaddr, BLKSECTS);
if (r < 0)
    panic("ide_read failed");
```

Finally, it remaps the page to clear the dirty bit:

```
r = sys_page_map(0, pgaddr, 0, pgaddr, PTE_W|PTE_U|PTE_P);
if (r < 0)
    panic("sys_page_map: %e", r);
```

The check that the block was allocated must be done after reading the block in, otherwise if the block is just the block that contains the free block bitmap, the check will cause a page fault again and finally leads to a triple fault.

In `flush_block()`, if the page is mapped and dirty, it writes a block out to disk using `ide_write()` and remaps the page to clear the dirty bit.

```
if (va_is_mapped(addr) && va_is_dirty(addr)) {
    void *pgaddr = ROUNDDOWN(addr, PGSIZE);
    r = ide_write(blockno * BLKSECTS, pgaddr, BLKSECTS);
    if (r < 0)
```

```

    panic("ide_write failed");
    r = sys_page_map(0, pgaddr, 0, pgaddr, PTE_W|PTE_U|PTE_P);
    if (r < 0)
        panic("sys_page_map: %e", r);
}

```

The Block Bitmap

Exercise 3

In `alloc_block()`, search the bitmap for a free block and mark it as in use. Immediately flush the changed bitmap block to disk.

```

int i;
for (i = 0; i < super->s_nblocks; ++i) {
    if (block_is_free(i)) {
        bitmap[i/32] &= ~(1<<(i%32));
        flush_block(bitmap + i/32);
        return i;
    }
}
return -E_NO_DISK;

```

File Operations

Exercise 4

In `file_block_walk()`, first check if `filebno` is out of range. Then if `filebno` is direct, the answer can be found in the `f_direct` array. Otherwise, we may need to allocate and initialize `f_indirect` block and find the answer there.

In `file_get_block()`, find the disk block number by `file_block_walk()`, allocate and initialize a block if it doesn't yet exist. Finally, set `*blk` to the address in memory where the `filebno`'th block of file 'f' would be mapped.

Client/Server File System Access

Exercise 5

In `serve_read()`, first look up the file id using `openfile_lookup()`. Then read the file content using `file_read()`. Finally, update the seek position.

In `devfile_read()`, first fill the `fsipcbuf` with the request arguments. Then make a request using `fsipc()`. Finally, move the result in `fsipcbuf` to the buffer.

Exercise 6

In `serve_write()`, first look up the file id using `openfile_lookup()`. Then write the content using `file_write()`. Finally, update the seek position.

In `devfile_write()`, first fill the `fsipcbuf` with the request arguments. Then make a request using `fsipc()`.

Client-Side File Operations

Exercise 7

In `open()`, first find an unused file descriptor page using `fd_alloc()`. Then fill the `fsipcbuf` with the request arguments, send the `FSREQ_OPEN` request and map the returned file descriptor page at the appropriate `fd` address with `fsipc()`. Finally, return the file descriptor index using `fd2num()`.

Spawning Processes

Exercise 8

In `sys_env_set_trapframe()`, first check that the `tf` pointer points to a good address. Then find the `Env` structure with `envid2env()` and copy the trap frame pointed by `tf`. Finally, modified the `tf` to make sure that user environments always run at CPL 3 with interrupts enabled.

For challenge, see the `answers-lab5.txt`.